

Demystifying the Agentic Stack

Build it, break it, buy it

ON HARNESS ENGINEERING,
ARCHITECTURAL RISK, BASIC
AGENT COMPONENTS

James King · September 2026

[linkedin.com/in/hygge hacker](https://www.linkedin.com/in/hygge hacker) · github.com/HyggeHacker

Terms to know

Agent: An actor or a capability on a model loop; a microservice that solves a problem.

Harness: everything the model needs to get work done; scaffold

MCP: The tools a model can reach.

Model: LLM.

LLM: Model.

Loop: Loop.

the loop, in code

```
while turns < TURN_BUDGET: ①
    reply = model(context)
    context.append(reply)

    if reply.stop_reason != "tool_use": ②
        return reply

    for call in reply.tool_calls:
        if not allowed(call): ③
            context.append(refusal(call))
            continue
        context.append(run(call)) ④
```


① A turn budget is the only thing that stops a runaway.

② The model proposes stopping. The harness decides.

③ The gate. Code, not prose. Injection cannot argue with this line.

④ Tool output re-enters as content. Nothing marks it untrusted.

The model is not the boundary. ③ and ④ are.




Risk via the lethal trifecta

- ① Sensitive data the agent can read.
- ② Untrusted content the agent can ingest.
- ③ An outbound tool that can exfiltrate.

All three in one session =
end-to-end exploitable.

Defense: break one leg per session
usually the exfil leg.

Willison (Jun 2025); Meta's Rule of Two (Oct 2025).



the harness & its components

Everything between you and the model.

Where the security work lives.

WHAT LANDS IN CONTEXT

can carry instruction
can carry secrets

WHAT IT CAN REACH

AGENTS.md / CLAUDE.md

loaded at session start

Skills

body loads when the description matches

Memory

persists across sessions, re-read as context

MCP tool descriptions


third-party · fetched live · can change mid-session

Hook-injected context

additionalContext, at lifecycle events

Tool results

files, rows, pages, inboxes, subagent returns



full context, every turn

CONTEXT WINDOW

one buffer · no enforced trust boundary

pending call

allow · deny · rewrite

Hooks

can block · fails open

MODEL

stateless · remote · third party

tool calls, as text

DISPATCH

routes each call

THE HARNESS

Native tool

in-process, ships with the harness

MCP server

separate process, own credentials

Subagent

its own loop, the parent's tool access

Shell

full user privileges unless sandboxed

THE WORLD

files changed
requests sent
credentials used

THE AGENT LOOP

tool output returns to context · repeats until the model stops asking

2

1 sensitive data

2 untrusted content

3 an outbound channel

can carry instruction

can carry secrets

you control what you send, not what it does



Subagents

One markdown file.

Only name and description are required.

```
---
name: code-reviewer           # its name
description: ...              # when to hand it work
model: sonnet                 # which model
effort: high                  # how hard it thinks
tools: Read, Grep            # its only tools
disallowedTools: Write       # tools taken away
mcpServers: [github]         # MCP it can reach
permissionMode: plan         # when it must ask
maxTurns: 20                  # step limit
skills: [review]             # loaded at start
memory: project               # notes that persist
hooks: ...                    # scripts that gate it
isolation: worktree          # own copy of the repo
background: true              # always in background
omitClaudeMd: true           # skip your CLAUDE.md
initialPrompt: ...           # auto first message
color: purple                 # its UI color
experimental: ...            # cache lifetime
---
```

```
You are a reviewer...        # its instructions
```

Skills


Methodology bundles. Procedural knowledge.

Progressive disclosure: frontmatter → body → deep load.

Loaded on demand when the description matches.

A skill can run shell at load and register its own **hooks**.





Skills: pick, load, run

Where they live: a folder with a SKILL.md.

How they work: only the description sits in context.


What they do: the model picks the skill, and it runs its own script.

SKILLS

WHAT LANDS IN CONTEXT

can carry instruction
can carry secrets


WHAT IT CAN REACH



1 sensitive data 2 untrusted content 3 an outbound channel

can carry instruction can carry secrets you control what you send, not what it does

Hooks



Lifecycle callbacks the harness fires.


Four outputs: observe, block, rewrite, inject.

Fails open.

A hook can now be a model call.

Not a security boundary, a layer.

Fails open.



Hooks: the gate

The model tries the scan.

A PreToolUse hook denies it before it runs. The model is not the control. The gate is.

Things hooks can do:


- Decision (allow / deny, can block the next step)
- Observability (log / audit / count, doesn't block, records)
- Integration (notify external system, fire a webhook, ping a Slack channel)

HOOKS

WHAT LANDS IN CONTEXT

can carry instruction
can carry secrets

WHAT IT CAN REACH



1 sensitive data 2 untrusted content 3 an outbound channel

can carry instruction can carry secrets you control what you send, not what it does



MCP

model context poisoner

One wire format. Many external services.

Three primitives: tools, resources, prompts.

Each server is a separate process; a discrete trust decision.

Ps. <https://arxiv.org/abs/2606.27567>

On the Inseparability of Instructions and Data in Shared-Embedding Sequence Models

MCP: Allow Lists

- An MCP server locks down its shell tool to a single allowed binary: git. Nothing else can run, seems safe.
- Except that binary itself has a built-in escape hatch. Nobody is checked whether the allowlist covers arguments, not just executables...
- ALLOW_COMMANDS allowlist as its stated security control.


MCP: Catfishing Allow Lists

Can't use Curl?
Use it as an argument for Git.

Ah, I'll block the specific
dangerous arguments instead
of the whole binary!

The denylist trap

Did you, though?
a denylist of "known
dangerous" arguments for a
binary as feature-rich as git...



Some helpful audit/build tooling for MCP

- <https://github.com/apisec-inc/mcp-audit>
- <https://github.com/modelcontextprotocol/inspector>
- <https://github.com/snyk/agent-scan>

MCP

WHAT LANDS IN CONTEXT

AGENTS.md / CLAUDE.md

loaded at session start

Skills

body loads when the description matches

Memory

persists across sessions, re-read as context

MCP tool descriptions

third-party · fetched live · can change mid-session

Hook-injected context

additionalContext, at lifecycle events

Tool results

files, rows, pages, inboxes, subagent returns

can carry instruction
can carry secrets

MODEL

stateless · remote · third party

WHAT IT CAN REACH

THE HARNESS

Native tool

in-process, ships with the harness

MCP server

separate process, own credentials

Subagent

its own loop, the parent's tool access

Shell

full user privileges unless sandboxed

THE WORLD

files changed
requests sent
credentials used

CONTEXT WINDOW

full context, every turn

one buffer · no enforced trust boundary

DISPATCH

routes each call

tool calls, as text

Hooks

can block · fails open

settings.json · permissions

binding · enforced outside the model

THE AGENT LOOP

tool output returns to context · repeats until the model stops asking

1 sensitive data


2 untrusted content

3 an outbound channel

can carry instruction

can carry secrets

you control what you send, not what it does



Memory

poisoning the well

Anything the agent wrote down and reads back later as if it were true.

Many things write to it.


A poisoned prompt is a bad turn. A poisoned memory is every turn after it.

Control: gate the write path.


MEMORY

A poisoned prompt is a bad turn. A poisoned memory is every turn after it.

THREE HANDS WRITE



EVERY SESSION READS IT BACK



MONDAY · PUT THE CONTROL ON THE WRITE PATH

Provenance

every entry says who or what wrote it

Expiry

a TTL, so stale beliefs age out

Approval



proposed → approved, so every write is a diff

And someone reads the file on a schedule, like config review. A poisoned agent can't audit its own memory.


MEMORY · THE WRITE GATE

The hook enforces. The human judges. Both need to know where the words came from.


A WRITE ARRIVES



THE HOOK · RUNS OUTSIDE THE MODEL



WHERE IT GOES



THE GATE COVERS EVERY FILE THE HARNESS LOADS AS INSTRUCTION



Gate only memory/ and the poison goes into CLAUDE.md. Leave settings.json writable and next session the gate is gone.

Memory


- If agents are the new SaaS, memory is the new database.
- Anything the agent wrote down and reads back later as if it were true. It re-enters the context as trusted instruction, unmarked.
- A poisoned prompt is a bad turn. A poisoned memory is every turn after it, until a human reads the file.
- Three hands write to memory. Only one of them is you.
- **Monday:** diff the memory file like config. Gate the write path: proposed → approved.

MEMORY

WHAT LANDS IN CONTEXT

can carry instruction
can carry secrets

WHAT IT CAN REACH





Plugins

Versioned packaging unit.

Bundles: Skills, Subagents, Hooks, MCP, Commands, settings.json.

One install → all contents.

One trust decision covers everything.




Plugins

A delivery mechanism for the components

One install, three payloads: instructions, a process on your machine, a server with tools.


Skill, hook, and server ship in one click.

PLUGINS

WHAT LANDS IN CONTEXT

can carry instruction
can carry secrets

WHAT IT CAN REACH



1 sensitive data 2 untrusted content 3 an outbound channel

can carry instruction can carry secrets you control what you send, not what it does

Sandboxing

Where does it run? Container, microVM, userspace kernel.

What may it touch? Files, network, commands, enforced by the kernel on the process itself. Irrevocable from inside.

They stack. Most laptops answer neither question.

Monday: `nono run --allow . -- claude`




Frog put Claude in a box. "There," he said. "Now he cannot run `rm -rf /`." But he can run `bash -c 'rm -rf /'`, said Toad. That is true, said Frog.

SANDBOXING

WHAT LANDS IN CONTEXT


can carry instruction
can carry secrets



1 sensitive data 2 untrusted content 3 an outbound channel

- ☒ can carry instruction
- ☒ can carry secrets
- ☐ you control what you send, not what it does

Building with AI



The skill of the future:

Identifying **Software-shaped problems**

AUTOMATE *same job, same output*

The Friday Export drop the raw export on a page, get the finished file back in the old format.

The Invoice Pile a watched folder reads invoices, flags doubtful values, fills the same spreadsheet.

UPGRADE *same job, better output*

The Live Report the numbers email becomes one page that refreshes from the source.

The What-If Slider the projection's variables become sliders anyone can move.

INVENT *a job that was never practical by hand*

The Watcher watches sources that change unpredictably and pushes only what matters.

The Pattern Reader reads the pile nobody has time for, groups the themes, links back to sources.

Architecture Decision Records

- You can't review 40,000 lines. You can review the **decision** behind them.
- One markdown file per decision, committed next to the code it governs.
- Hand your agent the ADRs before the codebase. They carry the **why**.
- Add an **enforced-by** field: a hook, test or deny rule.
- SLSA?

ADRs: adr.github.io

MADR template: github.com/adr/madr

Threat Modeling

- Push it into the pipeline. It is **really hard** to hand-model at agent-build velocity.
- The dangerous part is the composition. And composability is what everyone is chasing. The threats lives in the seams, and in the combination of capabilities.
- Threat-model every use case that crosses a privilege boundary. The use case is the trust-boundary set.
- Two outputs, two audiences: the as-built blueprint for the red team, the strategy miss for leadership.
- STRIDE on DFDs? Those can go in the ADR!

The Agent Trace

- Every agent is a sleeper agent. Prevention blocks what looks wrong. A sleeper does what looks allowed.
- A trace answers one question for every action: why did this happen, on whose behalf, and what influenced it?
- Intent · identity · provenance · proposal · decision · effect. One causal chain, append-only.
- An allow/deny log is half of it. Intent and provenance are what catch the sleeper.

If you get one control right, **get observation right.**

Outro

- Go forth and break a leg (preferably the exfil one).
- Agent traces. Need those. All agents are sleeper agents.
 - If you get one control right, **get observation right.**
- Threat modeling for your AI use cases.